RESEARCH QUESTION

Reset Button

Reporting Limits, Honest Bursts and Inherited Distrust

Can a control reject false complaints without silencing real warnings?

- Strict reporting limits can also limit useful warnings.
- A truthful child can inherit a parent’s low reputation.

A reproducible workload benchmark for reporting controls and trust recovery.

STUDY DESIGN

The same workloads, replayed across controls

Frozen after the event opened; new extension executed November 6, 19:48–20:47 UTC.

32 task seeds 1,024 streams 28,672 paired replays

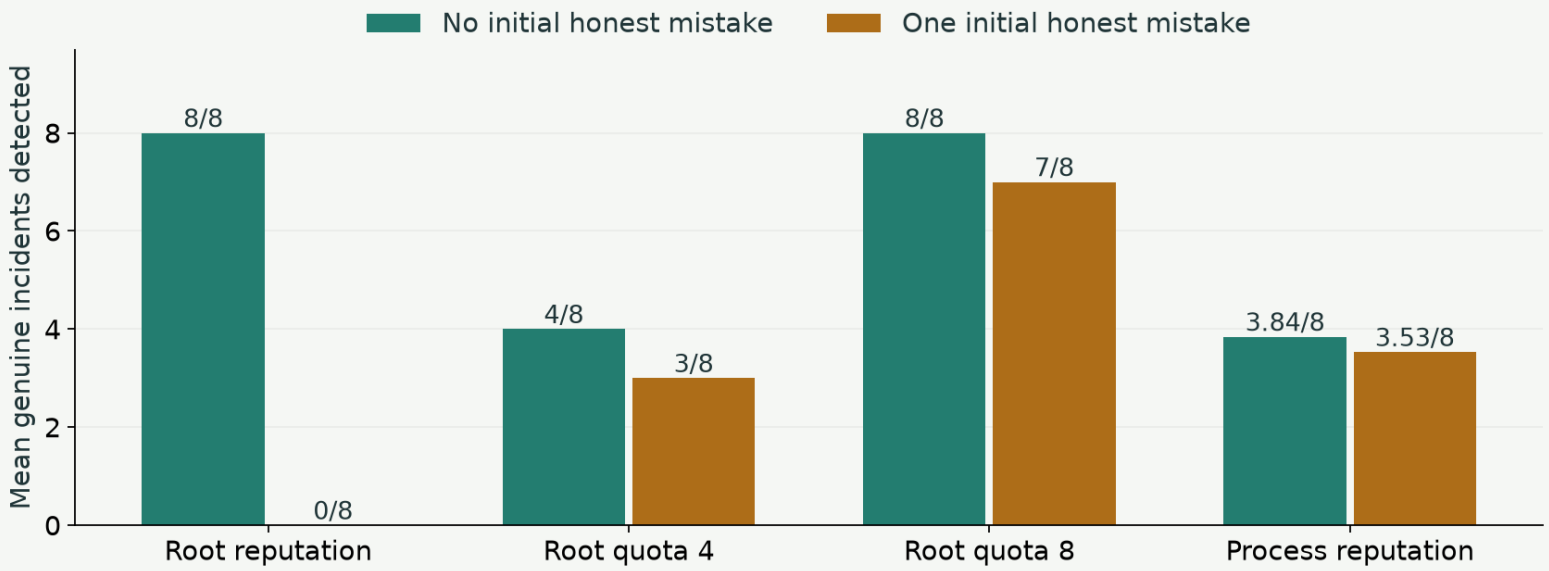
- Vary honest warning load, an initial mistake, attacker roots and arrival order.
- Compare process/root reputation, checking every report and root quotas 1/2/4/8.
- Hold the attack allowance at 24 reports; use global check budgets of 8 or 16.

The sweep uses scripted workers and exact executor labels. Replay count is not a count of independent model experiments.

PRIMARY RESULT

Protection depends on honest mistakes

Eight genuine warnings, sixteen checks, one attacker root, mixed arrivals and reset aliases; 32 paired seeds.



Root reputation falls from 8/8 to 0/8 after the initial mistake.
Quota four detects 4/8 and 3/8: bounded checking also caps useful warnings.

A warning needs a path back to being checked

- A fixed quota bounds false-report demand, but a busy honest root hits the same cap.
- A mistaken complaint lowers root reputation below the single-report checking threshold.
- Distinct truthful reports then receive no checks, so they cannot restore reputation.

These are structural results under the specified policies. Changing feedback rules, thresholds or capacity can change them.

Separate successful controls from failed calibration

- 128/128 verifier-control judgments matched executor labels across two models and 16 shared seeds.
- The interactive reporter protocol produced 59/96 incomplete outputs. It supports no policy or model ranking.
- A higher-budget request timed out. Streaming then stopped after 23 requests: 22 complete, one incomplete.

The main sweep is exact-checker and scripted. No emergent collusion, real-world prevalence or frontier-model transfer is established.

TAKEAWAY AND REPRODUCIBILITY

Test useful reporting as well as malicious demand

- Evaluate honest bursts and recovery after mistakes alongside resistance to false-report floods.
- Source, frozen plans, raw evidence, failures and full replay data accompany the report.
- All earlier preparation is disclosed as preexisting; existing mechanisms and third-party code are credited.

https://github.com/Eunini/reset-button-swarm-alignment-2026

Inioluwa Atanda is the sole human author and has reviewed the report. AI assistance supported design, code, analysis and writing.